

BiLSTM Model Documentation

1. Model Overview

The model is designed to predict **review ratings (1–5 stars)** from Amazon Fine Food Reviews text. It uses a **Bidirectional LSTM (BiLSTM)** architecture to capture the sequential and contextual information in the reviews.

Why BiLSTM:

- Reads text **forward and backward**, capturing context from both directions.
 - Handles long and complex sentences better than traditional models (like TF-IDF + Logistic Regression).
-

2. Data Preprocessing

- **Text Cleaning:** Removed punctuation, special characters, and converted text to lowercase.
 - **Tokenization:** Converted each review into a sequence of word indices.
 - **Padding Sequences:** Ensured all sequences have the same length for input to the BiLSTM.
 - **Derived Features:**
 - ReviewLength – number of words in each review
 - stopword_count – number of stopwords in each review
-

3. Model Architecture

1. **Embedding Layer** – Converts each word index into a dense vector of fixed size (captures semantic meaning).
2. **Bidirectional LSTM Layer** – Reads sequences in both forward and backward directions; includes dropout for regularization.
3. **Dense Layer(s)** – Fully connected layer(s) to output predicted rating.
4. **Output Layer** – Softmax activation for multi-class classification (ratings 1–5).

Hyperparameter Tuning

To optimize the BiLSTM model for the imbalanced Amazon Fine Food Reviews dataset, Keras Tuner was used with a Random Search strategy. The key goal was to maximize validation accuracy by exploring different combinations of important model hyperparameters.

Hyperparameters Tuned:

<u>Hyperparameter</u>	<u>Values Tested</u>	<u>Description</u>
Embedding Dimension	50, 100, 200	Size of the dense vector representing each word in the embedding layer
LSTM Units	64, 128	Number of hidden units in the BiLSTM layer
LSTM Dropout	0.2 – 0.5 (step 0.1)	Dropout applied to the BiLSTM layer to prevent overfitting
Dense Units	64, 128	Number of neurons in the fully connected dense layer
Dense Dropout	0.2 – 0.5 (step 0.1)	Dropout applied after the dense layer for regularization

Tuning Setup:

- RandomSearch was chosen to explore a random subset of the hyperparameter space efficiently.
- Max Trials: 3 (for faster tuning; can be increased for more exhaustive search)
- Executions per Trial: 1
- Objective: Maximize validation accuracy on the validation set.

Outcome / Notes:

- The tuner automatically evaluated different combinations of embedding size, LSTM units, and dropout rates.
- The best hyperparameters were selected based on the highest validation accuracy, balancing performance and model complexity.
- This tuned BiLSTM model was then used for final training and evaluation on the imbalanced dataset.

4. Training & Evaluation

- **Dataset Split:** Training, validation, and test sets
 - **Evaluation Metrics:** Accuracy, Precision, Recall, F1-score
 - **Observations:**
 - Training and validation loss decreased steadily.
 - The model performs better than classical ML models due to context capture in BiLSTM.
-

5. Tokenization & Embedding Summary

- **Tokenization:** Converts text into sequences of integers (word indices).
 - **Embedding Layer:** Maps each word index to a dense vector, allowing the model to learn **semantic relationships** between words.
-

6. Advantages

- Captures sequential context in reviews.
 - Handles longer texts with complex patterns.
 - More robust to imbalanced rating distribution than traditional ML models.
-

7. Final Notes

This BiLSTM model serves as a **production-ready text rating predictor** for Amazon Fine Food Reviews. It can be further optimized by tuning hyperparameters, adding more layers, or using pretrained embeddings (e.g., GloVe, Word2Vec).

Code:

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, Bidirectional, LSTM, Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping
import keras_tuner as kt
import tensorflow as tf

# 1. Load your dataset
# imbalanced_df has 'Review_text' and 'Rating'
```

```
X = imbalanced_df['Review_text'].astype(str).values
```

```
y = imbalanced_df['Rating'].values
```

```
# Encode labels
```

```
le = LabelEncoder()
```

```
y_encoded = le.fit_transform(y) # now labels: 0,1,2,3,4
```

```
# 2. Train-test split
```

```
X_train, X_test, y_train, y_test = train_test_split(
```

```
    X, y_encoded, test_size=0.2, stratify=y_encoded, random_state=42
```

```
)
```

```
# 3. Tokenization & Padding
```

```
max_words = 20000    # max words in vocabulary
```

```
max_len = 100        # max length of each review (reduces steps per epoch)
```

```
tokenizer = Tokenizer(num_words=max_words, oov_token="<OOV>")
```

```
tokenizer.fit_on_texts(X_train)
```

```
X_train_seq = tokenizer.texts_to_sequences(X_train)
```

```
X_test_seq = tokenizer.texts_to_sequences(X_test)
```

```
X_train_pad = pad_sequences(X_train_seq, maxlen=max_len, padding='post',  
truncating='post')
```

```
X_test_pad = pad_sequences(X_test_seq, maxlen=max_len, padding='post',  
truncating='post')
```

```
# 4. Build BiLSTM model function for hyperparameter tuning
```

```
def build_bilstm_model(hp):
```

```

model = Sequential()

model.add(Embedding(input_dim=max_words, output_dim=hp.Choice("embed_dim",
[50, 100, 200]), input_length=max_len))

model.add(Bidirectional(LSTM(units=hp.Choice("lstm_units", [64, 128]),
dropout=hp.Float("dropout", 0.2, 0.5, step=0.1))))

model.add(Dense(units=hp.Choice("dense_units", [64, 128]), activation='relu'))

model.add(Dropout(hp.Float("dense_dropout", 0.2, 0.5, step=0.1)))

model.add(Dense(5, activation='softmax')) # 5 classes


model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

return model

```

5. Hyperparameter Tuning

```

tuner = kt.RandomSearch(
    build_bilstm_model,
    objective='val_accuracy',
    max_trials=3,      # 3 trials for fast tuning
    executions_per_trial=1,
    directory='bilstm_tuning',
    project_name='imbalanced_reviews'
)

```

Early stopping

```

early_stop = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)

```

6. Run tuner

```
tuner.search(  
    X_train_pad, y_train,  
    validation_split=0.2,  
    epochs=15,          # keep 15 epochs  
    batch_size=128,     # increase batch size to reduce steps per epoch  
    callbacks=[early_stop],  
    verbose=1  
)
```

7. Get best model

```
best_model = tuner.get_best_models(num_models=1)[0]
```

Evaluate on test set

```
test_loss, test_acc = best_model.evaluate(X_test_pad, y_test, verbose=1)  
print("Test Accuracy:", test_acc)
```